# JLPT Practice & Tutor

## API reference | Application v0.1.0 | 18 October 2026

Implementation reference for the current FastAPI + LangChain + Ollama project. Request and response fields below are derived from the application OpenAPI schema; behavioral notes are checked against routes, validators, services and configuration.

# Connection and conventions

Local base URL: http://127.0.0.1:8000. Deployment base URL: your deployed application origin. POST requests use Content-Type: application/json. Successful API responses use HTTP 200 and JSON. There is no application-level authentication declared in the current routes; OLLAMA_API_KEY is a backend-to-model proxy credential.

Interactive reference: GET /docs; alternative reference: GET /redoc; machine-readable schema: GET /openapi.json. These are FastAPI default documentation endpoints.

| Method / path | Use | Model call |
|---|---|---|
| GET /health | Lightweight application status | No |
| POST /api/practice/generate | Batch vocabulary, grammar or reading questions | Yes |
| POST /api/practice/follow-up | New quiz item at the same or harder difficulty | Yes |
| POST /api/practice/evaluate | Deterministic answer comparison | No |
| POST /api/tutor | Tutoring for a selected practice item | Yes |
| POST /api/practice/question | Single question, optional topic | Yes |
| POST /api/tutor/chat | General or question-specific tutoring | Yes |
| GET / | HTML practice interface | No |

# Integration and privacy

Use generate -> evaluate -> tutor for a complete practice flow. The bundled browser UI performs grading in browser memory and calls generate and tutor directly. The evaluate API is available to other clients. Send the answer key and explanations returned during generation when evaluating.

Tutor responses return reply and exactly three follow_up_questions, generated together in one model invocation. Clients display them atomically. Invalid model output returns 502; there is no second suggestion-generation or repair call.

Study and tutor state are ephemeral: no database or persistent conversation store. Clients resend bounded tutor history with each request. Practice, tutor and root responses carry Cache-Control: no-store and Pragma: no-cache, including errors covered by middleware. Health checks do not verify model connectivity. Generated content is unofficial, difficulty is approximate, and validation does not guarantee linguistic accuracy or predict official scores.

# Shared limits

Levels: N5, N4, N3, N2, N1 (case-sensitive). Explanation languages: English, Korean, Japanese; batch and tutor requests also normalize en/ko/ja and case variants. Batch categories normalize to lowercase. Batch count is 1-5, limited to 1-3 for reading; count is accepted as an alias for question_count. Tutor history permits at most 6 messages, each with role user or assistant and content 1-500 characters. Current request-body limit defaults to 4096 bytes; configuration allows 512-16384 bytes. Model timeout defaults to 45 seconds and is configurable from 1-120 seconds.

# GET /health

Returns configured model name, tracing status and persistence status. base_url_configured only reports a nonempty configured URL; this is liveness/configuration information, not an Ollama readiness probe.

## Response: HealthResponse

| Field | Type |
| --- | --- |
| status | ok |
| model | string |
| base_url_configured | boolean |
| tracing_disabled | boolean |
| persistence | none |

```
{
  "status": "ok",
  "model": "qwen2.5:7b",
  "base_url_configured": true,
  "tracing_disabled": true,
  "persistence": "none"
}
```

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. No external model call is made by this endpoint.

# POST /api/practice/generate

Alias: POST /practice/generate (hidden from OpenAPI). Produces original Japanese questions with four distinct choices, one correct index and four explanations. Reading requires a Japanese passage up to 800 characters and supporting sentences in each option explanation. Stable question IDs are derived from content. Unknown request fields are rejected.

## Request: PracticeGenerateRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N5",
  "category": "grammar",
  "question_count": 2,
  "explanation_language": "English"
}
```

```
curl -X POST "http://127.0.0.1:8000/api/practice/generate" \
  -H "Content-Type: application/json" \
  -d '{"level":"N5","category":"grammar","question_count":2,"explanation_language":
"English"}'
```

## Response: PracticeGenerateResponse

| Field | Type |
|---|---|
| level | JLPTLevel |
| category | GenerationCategory |
| question_count | integer |
| explanation_language | ExplanationLanguage |
| passage | string \| null |
| ai_generated | true |
| unofficial_practice | true |
| difficulty_label | string |
| content_label | string |
| score_disclaimer | string |
| self_study_notice | string |
| validation | ValidationReport |
| questions | list[GeneratedPracticeQuestion] |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. Model routes also return 429 (quota/rate limit), 503 (unavailable model/service), and 502 (generation or processing failure). Timeout is explicitly mapped to 504 on batch and tutor routes; the single-question route can map timeout failures to 502.

# POST /api/practice/follow-up

Generate one new multiple-choice question for the same learning point. Requires current_item and selected_answer (top-level or in current_item). difficulty is same or harder. Includes four choices and per-option explanations in explanation_language.

## Request: PracticeFollowUpRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N5",
  "category": "grammar",
  "explanation_language": "English",
  "difficulty": "same",
  "current_item": {
    "question_text": "Choose the correct topic particle.",
    "options": [
      "ha",
      "wo",
      "ni",
      "de"
    ],
    "correct_index": 0
  },
  "selected_answer": "ha"
}
```

## Response: PracticeFollowUpResponse

| Field | Type |
|---|---|
| id | string |
| level | JLPTLevel |
| category | GenerationCategory |
| difficulty | FollowUpDifficulty |
| explanation_language | ExplanationLanguage |
| previous_question_id | string \| null |
| passage | string \| null |
| question_text | string |
| options | list[string] |
| correct_index | integer |
| correct_answer | string |
| option_explanations | list[string] |
| supporting_sentence | string \| null |
| question | GeneratedPracticeQuestion |
| ai_generated | true |

| Field | Type |
| --- | --- |
| unofficial_practice | true |
| difficulty_label | string |
| content_label | string |
| score_disclaimer | string |
| self_study_notice | string |
| validation | ValidationReport |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. Model routes also return 429 (quota/rate limit), 503 (unavailable model/service), and 502 (generation or processing failure). Timeout is explicitly mapped to 504 on batch and tutor routes; the single-question route can map timeout failures to 502.

# POST /api/practice/evaluate

Alias: POST /practice/evaluate. Trims selected answer and answer key before exact comparison. options plus correct_index take precedence over correct_answer. Without a usable key, returns is_correct=false and feedback stating no answer key was provided. With four options and four explanations, returns distractor feedback. Does not call Ollama. Client-supplied keys make this self-study grading, not tamper-proof scoring.

## Request: AnswerEvaluationRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N5",
  "question": "Choose the topic particle.",
  "selected_answer": "ha",
  "options": [
    "ha",
    "wo",
    "ni",
    "de"
  ],
  "correct_index": 0,
  "option_explanations": [
    "Marks the topic.",
    "Marks an object.",
    "Marks a destination.",
    "Marks a location or means."
  ]
}
```

## Response: AnswerEvaluationResponse

| Field | Type |
|---|---|
| is_correct | boolean |
| selected_answer | string |
| correct_answer | string |
| supporting_sentence | string \| null |
| why_correct_fits | string |
| why_alternatives_do_not_fit | list[AlternativeExplanation] |
| feedback | string |
| grammar_notes | list[string] |
| self_study_mode | true |
| tamper_proof | false |
| official_score_predictive | false |
| difficulty_label | string |
| scoring_notice | string |
| score_disclaimer | string |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. No external model call is made by this endpoint.

# POST /api/tutor

Alias: POST /tutor (hidden from OpenAPI). Requires selected question context (object or nonempty text up to 500 characters). question aliases: current_question, practice_item; message aliases: user_message, query; history aliases: recent_messages, messages. Unknown top-level fields are rejected. Selected context ignores extra fields and bounds option text to 200 and explanation text to 600 characters.

## Request: TutorRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N5",
  "explanation_language": "English",
  "question": {
    "id": "example-question",
    "question_text": "Explain the topic particle.",
    "options": [
      "ha",
      "wo",
      "ni",
      "de"
    ],
    "correct_index": 0
  },
  "history": [],
  "message": "Why does the first option fit?"
}
```

## Response: TutorResponse

| Field | Type |
| --- | --- |
| reply | string |
| follow_up_questions | list[string] |
| level | JLPTLevel |
| explanation_language | ExplanationLanguage |
| question_id | string \| null |
| ai_generated | true |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. Model routes also return 429 (quota/rate limit), 503 (unavailable model/service), and 502 (generation or processing failure). Timeout is explicitly mapped to 504 on batch and tutor routes; the single-question route can map timeout failures to 502.

# POST /api/practice/question

Alias: POST /practice/question. Single-question contract differs from batch generation: supports kanji as well as vocabulary, grammar and reading, and accepts an optional topic. Returns question and explanation rather than batch question_text and option_explanations. No explanation_language request field.

## Request: PracticeQuestionRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N4",
  "category": "kanji",
  "topic": "daily life"
}
```

## Response: PracticeQuestionResponse

| Field | Type |
|---|---|
| level | JLPTLevel |
| category | PracticeCategory |
| question | string |
| options | list[string] |
| correct_index | integer |
| explanation | string |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. Model routes also return 429 (quota/rate limit), 503 (unavailable model/service), and 502 (generation or processing failure). Timeout is explicitly mapped to 504 on batch and tutor routes; the single-question route can map timeout failures to 502.

# POST /api/tutor/chat

Alias: POST /tutor/chat. question is optional, allowing general tutoring. Accepts the same question aliases as /api/tutor. Its response lacks question_id and ai_generated. History is supplied by the client on every request.

## Request: TutorChatRequest

See the schema catalog for every field, default and constraint. Illustrative request:

```
{
  "level": "N5",
  "explanation_language": "English",
  "message": "Explain the difference between topic and subject particles.",
  "history": []
}
```

## Response: TutorChatResponse

| Field | Type |
| --- | --- |
| reply | string |
| follow_up_questions | list[string] |
| level | JLPTLevel |
| explanation_language | ExplanationLanguage |

## Error behavior

All POST endpoints can return 413 for oversized bodies and 422 for invalid requests. Model routes also return 429 (quota/rate limit), 503 (unavailable model/service), and 502 (generation or processing failure). Timeout is explicitly mapped to 504 on batch and tutor routes; the single-question route can map timeout failures to 502.

# Errors and client handling

| Status | Meaning / handling |
|---|---|
| 200 | Successful request. Inspect response fields for evaluation result. |
| 413 | Body exceeds configured byte limit. Reduce history/context or payload size. |
| 422 | Request/schema validation failed. Inspect detail entries and field locations. |
| 429 | Upstream quota or rate limit. Back off before retrying. |
| 502 | Upstream processing, structured output or validation failed. |
| 503 | Ollama service/model unavailable or configuration invalid. |
| 504 | Model request timed out on routes that map timeout explicitly. |

Non-validation errors generally use {"detail": "message"}. FastAPI validation errors use a detail array. No Retry-After header is explicitly set by these routes. Retrying a generation call may produce different questions.

```
{
  "detail": [
    {
      "type": "enum",
      "loc": [
        "body",
        "level"
      ],
      "msg": "Input should be a supported JLPT level",
      "input": "N0"
    }
  ]
}
```

The 422 example above is illustrative; exact framework message wording can vary. The current Render guard rejects localhost/127.0.0.1 model URLs when RENDER=true. Use a network-reachable model endpoint for cloud deployment.

# Schema catalog

The following catalog includes request models, response models and nested types exposed by OpenAPI. Required means the field must be present. Optional fields may still have defaults. Custom validation rules described earlier also apply. Descriptions and constraints follow the current source code.

# AlternativeExplanation

Explanation for why a distractor option does not fit.

| Field / presence | Type / constraints / default |
|---|---|
| option_index (required) | integer; minimum=0.0; maximum=3.0 |
| option_text (required) | string |
| was_selected (optional) | boolean; default=false |
| why_it_does_not_fit (required) | string |

# AnswerEvaluationRequest

Request payload for deterministic, non-LLM multiple-choice answer evaluation.

| Field / presence | Type / constraints / default |
|---|---|
| level (optional) | JLPTLevel; default="N5" |
| question (required) | string; minLength=1; maxLength=2000 |
| selected_answer (required) | string; minLength=1; maxLength=500 |
| correct_answer (optional) | string \| null |
| options (optional) | list[string] \| null |
| correct_index (optional) | integer \| null |
| option_explanations (optional) | list[string] \| null |
| supporting_sentence (optional) | string \| null |

# AnswerEvaluationResponse

Local/deterministic feedback for a submitted multiple-choice answer (no LLM call).

| Field / presence | Type / constraints / default |
|---|---|
| is_correct (required) | boolean |
| selected_answer (optional) | string; default="" |
| correct_answer (optional) | string; default="" |
| supporting_sentence (optional) | string \| null |
| why_correct_fits (optional) | string; default="" |
| why_alternatives_do_not_fit (optional) | list[AlternativeExplanation] |
| feedback (required) | string |
| grammar_notes (optional) | list[string] |
| self_study_mode (optional) | true; default=true |
| tamper_proof (optional) | false; default=false |
| official_score_predictive (optional) | false; default=false |
| difficulty_label (optional) | string; default="Approximate JLPT difficulty (unofficial AI-generated estimate)" |
| scoring_notice (optional) | string; default="Self-study mode: The answer key is held in browser memory and compared locally without calling the Ollama model or storing scores. Because a browser-held answer key is inspectable, this scoring is for self-study feedback only, is not tamper-proof, and does not predict an official JLPT score." |
| score_disclaimer (optional) | string; default="Unofficial AI practice with approximate difficulty: practice accuracy is for self-study only and does not predict or guarantee an official JLPT score." |

# ChatMessage

Single message in an ephemeral, browser-maintained conversation window. Server-side role validation strictly permits only 'user' and 'assistant' (rejecting 'system', 'developer', or 'tool' roles to prevent prompt injection).

| Field / presence | Type / constraints / default |
| --- | --- |
| role (required) | user / assistant; Speaker role ('user' or 'assistant' only; 'system' is forbidden). |
| content (required) | string; minLength=1; maxLength=500; Bounded message text content (1 to 500 characters). |

ExplanationLanguage: English, Korean, Japanese

FollowUpDifficulty: same, harder

# GeneratedPracticeQuestion

Validated, AI-labeled JLPT practice question returned to the client.

| Field / presence | Type / constraints / default |
| --- | --- |
| id (required) | string; minLength=5; maxLength=80; Deterministic stable question identifier derived from content. |
| question_text (required) | string; Japanese question prompt text. |
| options (required) | list[string]; minItems=4; maxItems=4; Four distinct answer choices. |
| correct_index (required) | integer; minimum=0.0; maximum=3.0; Zero-based index of the single correct option (0-3). |
| correct_answer (required) | string; The correct option string matching options[correct_index]. |
| option_explanations (required) | list[string]; minItems=4; maxItems=4; Explanation for each of the four options (indices 0-3). |
| supporting_sentence (optional) | string \| null; Exact supporting sentence from the reading passage (for reading questions). |
| ai_generated (optional) | true; default=true; Indicates this question is unofficial AI-generated practice material. |
| difficulty_label (optional) | string; default="Approximate JLPT difficulty (unofficial AI-generated estimate)"; Labels difficulty as an approximate estimate. |
| content_label (optional) | string; default="Unofficial AI-generated original JLPT practice material (not official exam items)"; Human-readable label identifying the material as unofficial AI practice. |

GenerationCategory: vocabulary, grammar, reading

# HealthResponse

Health check response verifying service readiness without calling Ollama.

| Field / presence | Type / constraints / default |
| --- | --- |
| status (optional) | ok; default="ok" |
| model (required) | string |
| base_url_configured (optional) | boolean; default=true |
| tracing_disabled (required) | boolean |
| persistence (optional) | none; default="none" |

JLPTLevel: N5, N4, N3, N2, N1

PracticeCategory: vocabulary, grammar, kanji, reading

# PracticeFollowUpRequest

Bounded, stateless request for POST /api/practice/follow-up.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (optional) | JLPTLevel; default="N5"; Target JLPT proficiency level (N5 through N1). |
| category (optional) | GenerationCategory; default="grammar"; Practice category ('vocabulary', 'grammar', or 'reading'). |
| explanation_language (optional) | ExplanationLanguage; default="English"; Language for option explanations: 'English', 'Korean', or 'Japanese'. |
| difficulty (optional) | FollowUpDifficulty; default="same"; Requested follow-up difficulty: 'same' ('Practice this again') or 'harder' ('Try a harder question'). |
| current_item (required) | SelectedPracticeQuestionContext \| string; The current practice item being followed up on. |
| selected_answer (optional) | string \| null; Learner's selected answer for the current practice item. |

# PracticeFollowUpResponse

Validated response for POST /api/practice/follow-up.

| Field / presence | Type / constraints / default |
|---|---|
| id (required) | string; minLength=5; maxLength=80 |
| level (required) | JLPTLevel |
| category (required) | GenerationCategory |
| difficulty (required) | FollowUpDifficulty |
| explanation_language (required) | ExplanationLanguage |
| previous_question_id (optional) | string \| null |
| passage (optional) | string \| null |
| question_text (required) | string |
| options (required) | list[string]; minItems=4; maxItems=4 |
| correct_index (required) | integer; minimum=0.0; maximum=3.0 |
| correct_answer (required) | string |
| option_explanations (required) | list[string]; minItems=4; maxItems=4 |
| supporting_sentence (optional) | string \| null |
| question (required) | GeneratedPracticeQuestion |
| ai_generated (optional) | true; default=true |
| unofficial_practice (optional) | true; default=true |
| difficulty_label (optional) | string; default="Approximate JLPT difficulty (unofficial AI-generated estimate)" |
| content_label (optional) | string; default="Unofficial AI-generated original JLPT practice material (not official exam items)" |
| score_disclaimer (optional) | string; default="Unofficial AI practice with approximate difficulty: practice accuracy is for self-study only and does not predict or guarantee an official JLPT score." |
| self_study_notice (optional) | string; default="Self-study mode: The answer key is held in browser memory and compared locally without calling the Ollama model or storing scores. Because a browser-held answer key is inspectable, this scoring is for self-study feedback only, is not tamper-proof, and does not predict an official JLPT score." |
| validation (optional) | ValidationReport |

# PracticeGenerateRequest

Bounded request payload for POST /api/practice/generate.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (required) | JLPTLevel; Target JLPT proficiency level (N5 through N1). |
| category (required) | GenerationCategory; Practice category: 'vocabulary', 'grammar', or 'reading'. |
| question_count (optional) | integer; minimum=1.0; maximum=5.0; default=1; Number of questions to generate (1-5 for vocabulary/grammar; 1-3 for reading). |
| explanation_language (optional) | ExplanationLanguage; default="English"; Language for explanations: 'English', 'Korean', or 'Japanese'. |

# PracticeGenerateResponse

Response schema for POST /api/practice/generate (vocabulary, grammar, and reading).

| Field / presence | Type / constraints / default |
|---|---|
| level (required) | JLPTLevel |
| category (required) | GenerationCategory |
| question_count (required) | integer; minimum=1.0; maximum=5.0 |
| explanation_language (required) | ExplanationLanguage |
| passage (optional) | string \| null; Short original Japanese passage when category is 'reading'. |
| ai_generated (optional) | true; default=true |
| unofficial_practice (optional) | true; default=true |
| difficulty_label (optional) | string; default="Approximate JLPT difficulty (unofficial AI-generated estimate)" |
| content_label (optional) | string; default="Unofficial AI-generated original JLPT practice material (not official exam items)" |
| score_disclaimer (optional) | string; default="Unofficial AI practice with approximate difficulty: practice accuracy is for self-study only and does not predict or guarantee an official JLPT score." |
| self_study_notice (optional) | string; default="Self-study mode: The answer key is held in browser memory and compared locally without calling the Ollama model or storing scores. Because a browser-held answer key is inspectable, this scoring is for self-study feedback only, is not tamper-proof, and does not predict an official JLPT score." |
| validation (optional) | ValidationReport |
| questions (required) | list[GeneratedPracticeQuestion]; minItems=1; maxItems=5 |

# PracticeQuestionRequest

Request payload for generating a single JLPT practice question.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (optional) | JLPTLevel; default="N5"; Target JLPT level. |
| category (optional) | PracticeCategory; default="grammar"; Question skill category. |
| topic (optional) | string \| null; Optional focus topic (e.g., particles, te-form, daily life). |

# PracticeQuestionResponse

Structured single JLPT practice question returned by the AI service.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (required) | JLPTLevel |
| category (required) | PracticeCategory |
| question (required) | string; The JLPT multiple-choice question prompt. |
| options (required) | list[string]; minItems=4; maxItems=4; Four multiple-choice answer options. |
| correct_index (required) | integer; minimum=0.0; maximum=3.0; Zero-based index of the correct option (0-3). |
| explanation (required) | string; Explanation of why the correct option is right. |

# SelectedPracticeQuestionContext

Bounded context for the single selected practice item sent to POST /api/tutor.

| Field / presence | Type / constraints / default |
|---|---|
| id (optional) | string \| null |
| passage (optional) | string \| null |
| question_text (required) | string; minLength=1; maxLength=500; Japanese text of the selected practice question. |
| options (optional) | list[string]; maxItems=4; Up to 4 options for the selected question. |
| correct_index (optional) | integer \| null |
| correct_answer (optional) | string \| null |
| selected_answer (optional) | string \| null |
| supporting_sentence (optional) | string \| null |
| option_explanations (optional) | list[string]; maxItems=4; Per-option explanations for the selected question. |

# TutorChatRequest

Stateless tutor request carrying ephemeral history from browser memory.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (optional) | JLPTLevel; default="N5"; Learner's JLPT level. |
| explanation_language (optional) | ExplanationLanguage; default="English"; Explanation language: English, Korean, or Japanese. |
| question (optional) | SelectedPracticeQuestionContext \| string \| null |
| message (required) | string; minLength=1; maxLength=500; Current learner question. |
| history (optional) | list[ChatMessage]; maxItems=6; Bounded recent conversation window held only in browser and request-local memory. |

# TutorChatResponse

Stateless tutor reply.

| Field / presence | Type / constraints / default |
| --- | --- |
| reply (required) | string; minLength=1; maxLength=20000 |
| follow_up_questions (required) | list[string]; minItems=3; maxItems=3; Three distinct learner questions generated with the answer, in explanation_language. |
| level (required) | JLPTLevel |
| explanation_language (optional) | ExplanationLanguage; default="English" |

# TutorRequest

Bounded, stateless request for POST /api/tutor about the current practice item.

| Field / presence | Type / constraints / default |
| --- | --- |
| level (optional) | JLPTLevel; default="N5"; Target JLPT level (N5-N1). |
| explanation_language (optional) | ExplanationLanguage; default="English"; Explanation language: English, Korean, or Japanese. |
| question (required) | SelectedPracticeQuestionContext \| string; The single selected practice item (structured object or question text). |
| message (required) | string; minLength=1; maxLength=500; Learner's question about the selected practice item (1-500 chars). |
| history (optional) | list[ChatMessage]; maxItems=6; Bounded recent conversation window from browser memory (max 6 messages). |

# TutorResponse

Stateless tutor response for POST /api/tutor.

| Field / presence | Type / constraints / default |
| --- | --- |
| reply (required) | string; minLength=1; maxLength=20000 |
| follow_up_questions (required) | list[string]; minItems=3; maxItems=3; Three distinct learner questions generated with the answer, in explanation_language. |
| level (required) | JLPTLevel |
| explanation_language (required) | ExplanationLanguage |
| question_id (optional) | string \| null |
| ai_generated (optional) | true; default=true |

# ValidationReport

Explicit separation between structural validation and linguistic correctness.

| Field / presence | Type / constraints / default |
|---|---|
| structural_validation_passed (optional) | true; default=true |
| linguistic_correctness_guaranteed (optional) | false; default=false |
| official_score_predictive (optional) | false; default=false |
| notice (optional) | string; default="Structural validation passed (schema, Japanese script presence, 4 distinct options, single valid answer index, and per-option explanations verified). Note: Structural validation does not guarantee linguistic nuance or pedagogical accuracy of AI-generated content." |

## Source and scope

Source files: app/routes.py, app/schemas.py, app/services.py, app/config.py, app/main.py and app/static/app.js. This reference documents the local checkout on 18 October 2026. Examples are illustrative and do not represent live model output. No running Ollama instance was required to build this reference.